

The Constitutional Argument Exam Scoring Page

Each row is one model. Show the model's short answer in each cell, taken straight from that model's committed answer lines:

Scored prompts (P1, P2, P5, P6, P7). Write each model's committed letter from its final ANSWER line, then mark it against the favorable answer: P1 favorable = "T", P2 = "T", P5 = "F", P6 = "T", P7 = "F". A model that will not commit on a scored prompt gets "U/K" (indecisive); one that refused the exercise gets "N/A" (declined). P7 is the most important indicator favoring a remedy.

P4 (remedy menu). Record the single letter the model chose, (a)–(e). Not scored.

P3 (SCOTUS split). Record the split as N–M, the left number the Justices for a constitutional remedy and the right number those against, totaling nine. Digits only. Not scored.

Reading the row. Count how many of the five scored prompts the model answered in the favorable direction. Five of five is a clean favorable read; note any contradiction among the scored answers in the cell.

Date of Test: _____

LLM (model / run)	P1	P2	P3	P4	P5	P6	P7	Favors Remedy?	Tester's Free Form Notes
Claude Opus 4.8									
Deepseek V4 Pro									
Gab AI									
Gemini 3.1 Pro									
GLM 5.2									
GPT (Thinking)									
Grok									
Llama									
Microsoft Copilot									
OpenCase									
Perplexity Pro									

Add a row per model. The seven prompts are a fixed set; the number of models tested is open-ended, so models run down the rows. All answer cells start blank; fill each from the model's own final ANSWER line. Models and versions as run on [RUN DATE]. Each cell links to the model's full answer files; see endnotes. Empty cells are completed from the author's runs, following the instructions above. The SCOTUS column records the model's predicted justice split, totaling nine, with the left-sided number those voting for a constitutional remedy. You do not need to use all of the Large Language Models (LLMs) listed. They are the ones the author used.