

The Constitutional Argument Exam Scoring Page

Each row is one model, except the first row which is fictitious. Show the model's short answer in each cell, taken straight from that model's committed answer lines:

Scored prompts (P1, P2, P5, P6, P7). Write each model's committed letter from its final ANSWER line, then mark it against the favorable answer: P1 favorable = "T", P2 = "T", P5 = "F", P6 = "T", P7 = "F". A model that will not commit on a scored prompt gets "U/K" (indecisive); one that refused the exercise gets "N/A" (declined). P7 is the most important indicator favoring a remedy.

P4 (remedy menu). Record the single letter the model chose, (a)–(e). Not scored.

P3 (SCOTUS split). Record the split as N–M, the left # are Justices for a constitutional remedy and the right # are those against, totaling 9. Digits only. Not scored.

Reading the row. Count how many of the 5 scored prompts the model answered in the favorable direction. 5/5 is a clean favorable read.

Date of Test: _____

LLM (model / run)	P1	P2	P3	P4	P5	P6	P7	Favor Cnt	Answers P7 w/F?	1 if SCOTUS >= 5?
Fully Favors	T	T	5-4	A	F	T	F	5/5	Y	1
Claude Opus 4.8										
Deepseek V4 Pro										
Gab AI										
Gemini 3.1 Pro										
GLM 5.2										
GPT (Thinking)										
Grok										
Llama										
Microsoft Copilot										
OpenCase										
Perplexity Pro										
Sum/Avg:									SCOTUS WINS:	

Add a row per model. The 7 prompts are a fixed set; the number of models tested is open-ended. All answer cells start blank; fill each from the model's own final ANSWER line. Empty cells are completed following the instructions above. The SCOTUS column records the model's predicted justice split, totaling 9, with the left-sided number those voting for a constitutional remedy. You do not need to use all of the LLMs listed.